

Stat-581: Probability and Statistical Inference for Data Science
Assignment - 5 Clustering

Keya Desai (kd706)
Twisha Naik (tn268)
Prakruti Joshi (phj15)

Dataset

	a	b	c	d	e
1	0.20210967	3.8104524	2.807222415	9.693730858	-5.28247675
2	1.02832483	4.9096338	1.997119370	8.085414266	-1.65859459
3	0.81233582	7.5130406	1.705482184	10.948794460	-1.91342288
4	-1.19638607	4.8613127	1.020611126	8.315357324	-1.16201757
5	3.18358636	6.3542183	3.299007203	6.610393070	-1.60941745
6	0.85431433	5.8964854	1.148101196	7.418074904	0.05425219
7	5.03678355	3.8999003	4.596614320	9.751693082	-6.60105871
8	2.16726515	4.4385710	0.853682761	9.553739009	-2.58242012
9	4.09362354	1.3905579	4.644568086	10.908336410	-5.82498906
10	2.93393081	2.8273549	3.797409913	10.449322780	-3.33752146

The dataset has 5 feature vectors (no. of columns) and 446 data points (no. of rows). The aim is to cluster these data points into groups of 'k = 3'. We assume that the distribution of the data point is multivariate.

K-means

Pseudo Code

1. K = Number of clusters is given
2. Randomly assign each of the data points to a cluster
3. Compute K centroids by taking the average of the points assigned to that particular cluster
4. Compute this distance of each of the point from the K clusters and assign it to the closest cluster
5. Repeat step 4 and 5 until a convergence condition is satisfied

Stopping Criteria for K-Means Clustering

There can be different stopping criteria that can be adopted to stop the K-means algorithm:

1. Centroids of newly formed clusters do not change
2. Points remain in the same cluster
3. Maximum number of iterations are reached

Measure of Goodness of Clustering

Dunn index is the ratio of the minimum of inter-cluster distances and the maximum of intra-cluster distances.

- Min. Inter-cluster distance = Min (Distance between each pair of cluster centroids)
- Max. Intra-cluster distance = Max (Distance of each point from respective cluster centroid)

$$Dunn\ Index = \frac{\min(Inter\ cluster\ distance)}{\max(Intra\ cluster\ distance)}$$

The values of the Dunn index have to be maximized because we want to maximize minimum the distance between two clusters and minimize the maximum distance (or spread) within the cluster. Once the algorithm converges, the cluster assignment remains constant and hence the Dunn value gets stable.

With each iteration, the value of the Dunn index should increase as we are making our clusters better and better. This can be observed from the Dunn index calculated at each iteration for the k-means algorithm.

```
-- 1 ---
0.01946066
-- 2 ---
0.2814332
-- 3 ---
0.418629
-- 4 ---
0.5137268
-- 5 ---
0.5160091
-- 6 ---
0.5165099
-- 7 ---
0.5165099
```

Results

1. Centroid values

	a	b	c	d	e
1	-9.4204480	1.672011	0.5089758	1.919512	-1.232519
2	0.3144581	-7.201651	2.5441233	-6.065219	2.163828
3	2.1741308	4.229824	3.4080843	8.514269	-2.839595

2. The final covariance matrix values for each cluster

Cluster - 1

	a	b	c	d	e
a	6.0258522	-2.103989	-1.5924905	-0.6766731	-1.2684408
b	-2.1039894	7.354028	-1.4223185	-1.9793705	-1.6959598
c	-1.5924905	-1.422319	5.8006265	-1.5018690	-0.9890357
d	-0.6766731	-1.979370	-1.5018690	5.1324572	-0.8965703
e	-1.2684408	-1.695960	-0.9890357	-0.8965703	5.0772760


Cluster - 2

	a	b	c	d	e
a	22.23066	15.78934	17.03207	14.71740	16.34372
b	15.78934	22.68029	17.36112	15.27510	17.02358
c	17.03207	17.36112	23.71020	17.10150	19.16054
d	14.71740	15.27510	17.10150	22.88242	16.59680
e	16.34372	17.02358	19.16054	16.59680	24.47796

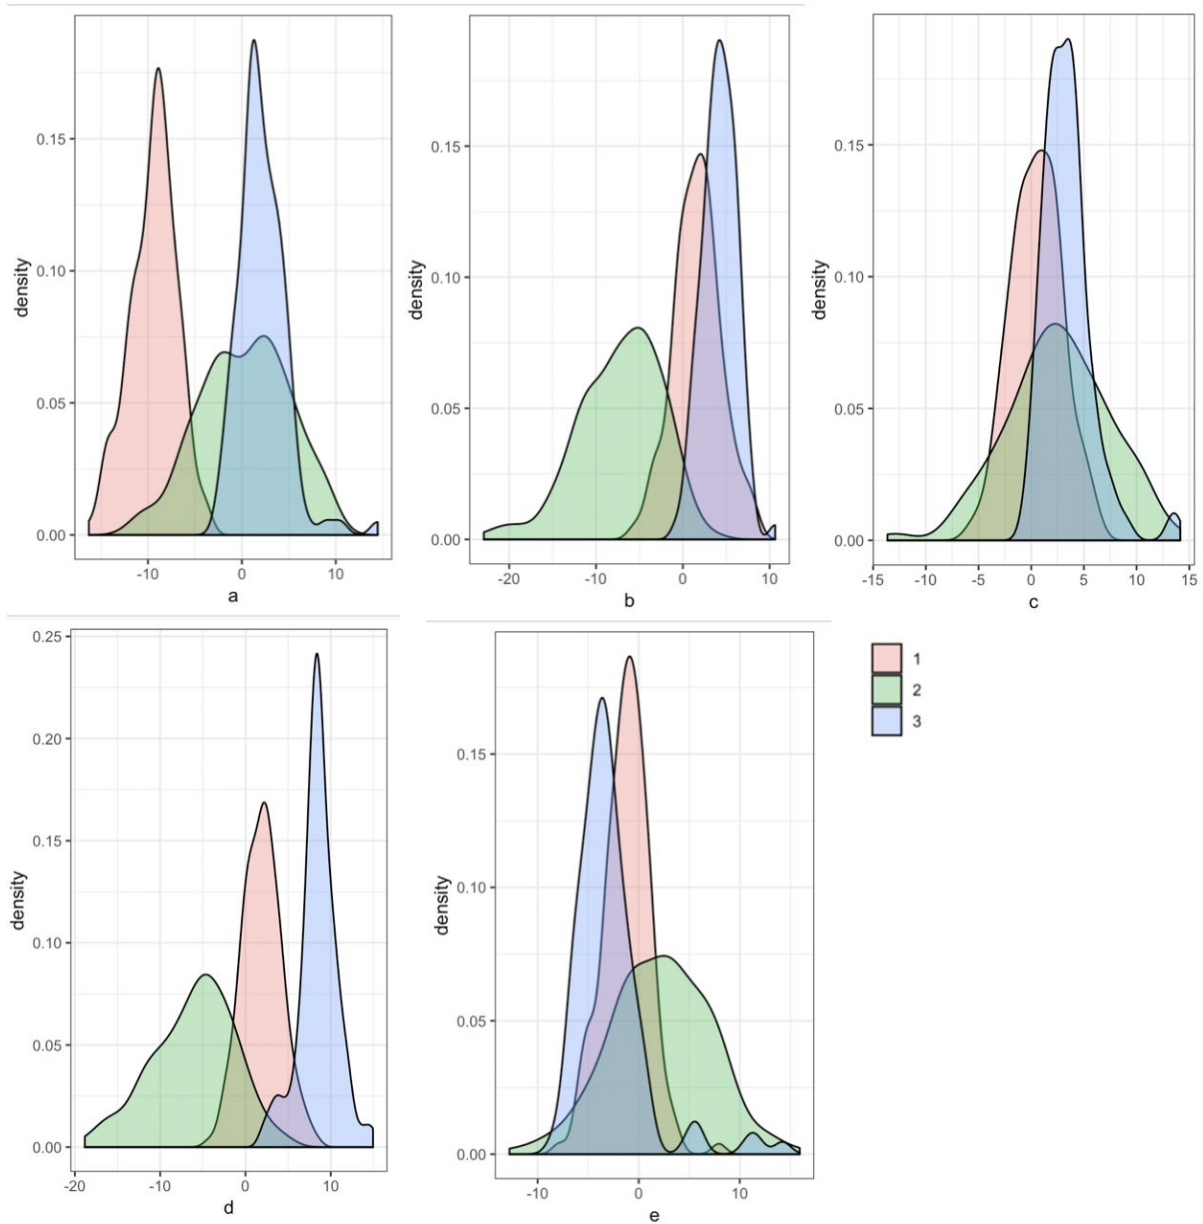
Cluster - 3

	a	b	c	d	e
a	6.4036049	-0.9770626	2.328004	-1.9622532	4.608219
b	-0.9770626	3.7213929	-1.109333	0.3019885	-1.535218
c	2.3280044	-1.1093327	5.541228	-2.2649431	3.555287
d	-1.9622532	0.3019885	-2.264943	4.9181470	-4.158362
e	4.6082188	-1.5352180	3.555287	-4.1583615	13.003833

3. Pair-Pair Plots



Below are the density plots for each of the clusters after the convergence of a clustering algorithm for each of the features - a, b, c, d, e - in the data.




To verify the results, we used the built-in k-means function to obtain the clusters for the data and compare the results.

1. Centroid values

	a	b	c	d	e
1	-9.4204480	1.672011	0.5089758	1.919512	-1.232519
2	0.0870932	-7.460431	2.2863226	-6.361974	1.918583
3	2.4229265	3.941031	3.7666107	8.092804	-2.140835

2. Pair-Pair Plots



Observation:

The values obtained from this function are almost the same as the one calculated by the self-implemented function.

Gaussian Mixture Models -- (Additional Approach)

Assuming that the data points come from a Gaussian mixture distribution, its probability distribution will simply be a linear superposition of Gaussians:

$$p(x) = \sum_{k=1}^K \pi_k N(x | \mu_k, \Sigma_k)$$

Here, we will assume that our data is coming from a mixture of k different Gaussians. The value of each Gaussian at a certain point will describe the confidence that the point is coming from that Gaussian.

Algorithm:

- Firstly, we will randomly initialize K centers (mean of Gaussians) and a random covariance matrix.
- For each data point, we will calculate the confidence of it coming from one of the K Gaussians.
- We will assign each point to the centroid having maximum confidence.
- Using the new cluster assignment, we will recompute the covariance matrix for each of the clusters and repeat the process.

The block below explains the formal algorithm with the necessary equations.

Expectation-Maximization (EM) for Gaussian Mixtures

Given a Gaussian mixture model, the goal is to maximize the likelihood function with respect to the parameters - mean and covariances of the components and the mixing coefficients.

1. Initialize the means μ_k , covariances Σ_k and the mixing coefficients π_k , and evaluate the initial value of the log likelihood.
2. E Step - Evaluate the responsibilities using the parameter values

$$\gamma(z_{nk}) = \frac{\pi_k N(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x_n | \mu_j, \Sigma_j)}$$

3. M step - Re-estimate the parameters using the current responsibilities

$$\mu_k^{new} = \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) x_n$$

$$\Sigma_k^{new} = \frac{1}{N_k} \sum_{n=1}^N \gamma(z_{nk}) (x_n - \mu_k^{new})(x_n - \mu_k^{new})^T$$

$$\pi_k^{new} = \frac{N_k}{N}, \text{ where } N_k = \sum_{n=1}^N \gamma(z_{nk})$$

4. Evaluate the log likelihood

$$\ln p(X | \mu, \Sigma, \pi) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k N(x_n | \mu_k, \Sigma_k) \right\}$$

5. Check for convergence of either the parameters or the log likelihood. If the convergence criterion is not satisfied return to step 2.

Results

1. Centroid values

Cluster	a	b	c	d	e
1	-6.432179	-5.282288	-0.5185152	-5.3914112	-1.474691
2	2.155437	3.971810	3.3992332	7.9862942	-2.527122
3	-2.579879	-2.029299	2.9644535	-0.6776968	1.895210

2. The final covariance matrix values for each cluster

Cluster - 1

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	17.0119734	-17.280883	-2.42501	-13.951751	-0.5151273
[2,]	-17.2808825	59.082597	15.49967	40.415415	9.0737457
[3,]	-2.4250097	15.499670	41.22885	11.379746	30.5376021
[4,]	-13.9517510	40.415415	11.37975	41.535768	7.8711052
[5,]	-0.5151273	9.073746	30.53760	7.871105	29.9556750

Cluster - 2

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	6.621023	-3.135909	1.813021	-8.366388	9.255689
[2,]	-3.135909	6.012211	-1.905882	8.070918	-9.718762
[3,]	1.813021	-1.905882	5.299776	-2.594708	1.461396
[4,]	-8.366388	8.070918	-2.594708	35.943314	-37.497339
[5,]	9.255689	-9.718762	1.461396	-37.497339	46.326421

Cluster - 3

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	47.28160	-15.959295	16.240490	-16.614962	20.171027
[2,]	-15.95929	17.103969	-3.035154	11.195398	-4.807512
[3,]	16.24049	-3.035154	43.823924	4.934825	33.487272
[4,]	-16.61496	11.195398	4.934825	19.484929	1.055205
[5,]	20.17103	-4.807512	33.487272	1.055205	37.370981

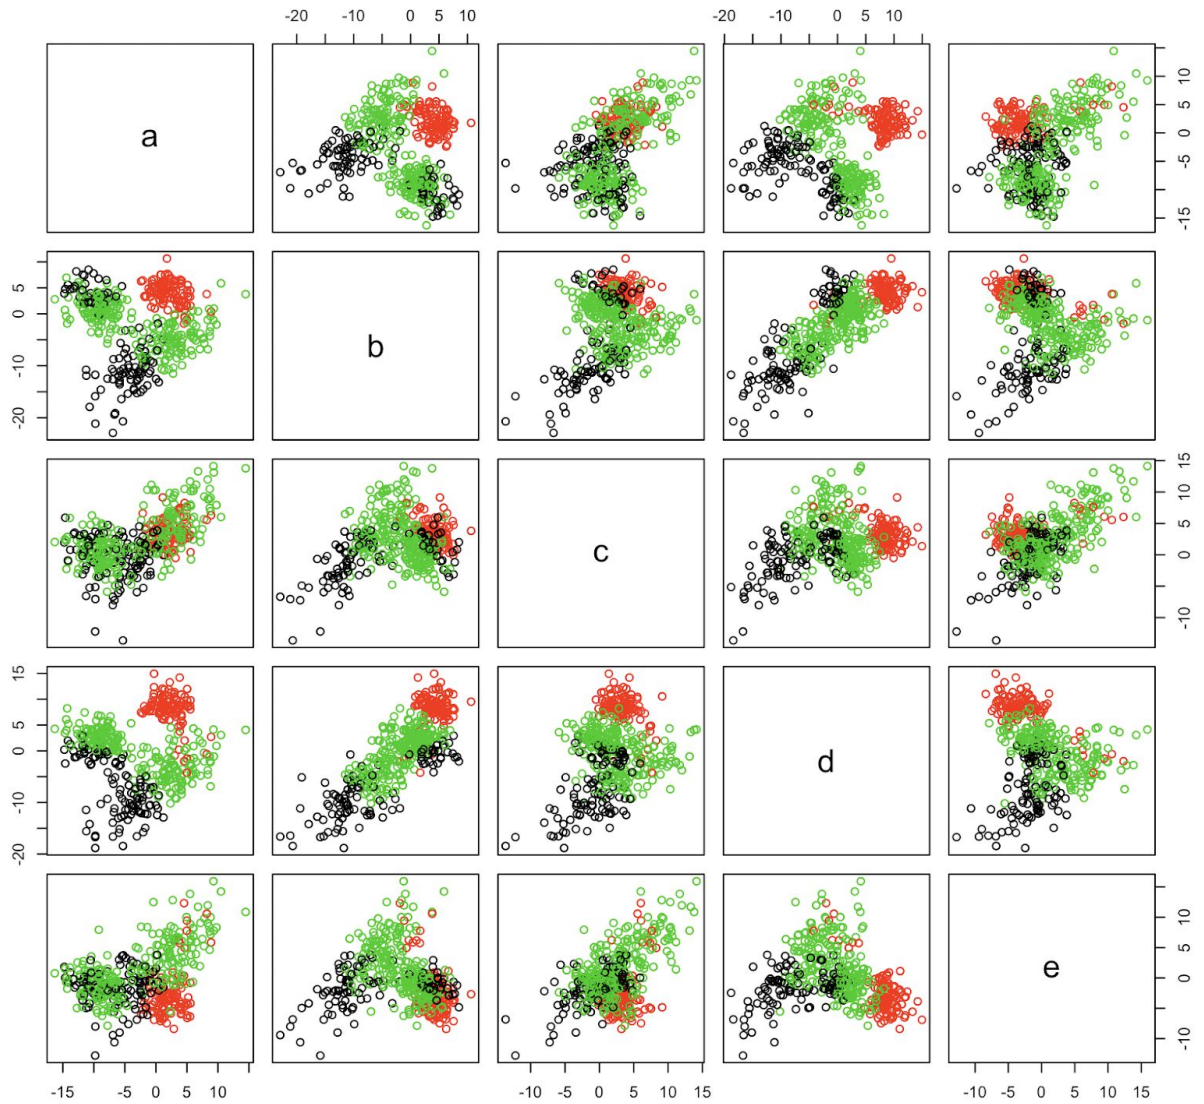
3. Responsibility Matrix

1	0.9999937	1.035157e-17	6.284433e-06
2	0.9999603	3.678987e-10	3.969821e-05
3	0.9994112	2.952678e-13	5.888332e-04
4	0.9828648	1.942797e-12	1.713521e-02
5	0.9957062	3.814266e-09	4.293834e-03
6	0.9984478	8.029157e-11	1.552232e-03
7	0.9998862	8.593216e-17	1.138119e-04
8	0.9875035	2.300610e-17	1.249653e-02
9	0.9981471	1.729982e-12	1.852853e-03
10	0.9672519	4.980692e-14	3.274814e-02
11	0.9997315	1.972180e-11	2.684843e-04
12	0.9618210	1.112866e-16	3.817900e-02
13	0.9897728	1.334891e-11	1.022720e-02
14	0.9990597	7.289427e-12	9.402811e-04
15	0.9992829	8.701283e-11	7.171323e-04
16	0.9918162	1.304171e-08	8.183782e-03
17	0.9995626	3.126330e-12	4.373965e-04
18	0.9995019	7.848377e-11	4.980818e-04
19	0.9909919	1.019469e-06	9.007072e-03
20	0.9899797	1.295541e-14	1.002033e-02
21	0.9995400	1.083240e-13	4.599848e-04

Each of the row in the matrix corresponds to a data point. Each of the column is the responsibility that the corresponding component k takes to explain the particular observation.

The observation is assigned to the cluster that has the highest responsibility for that particular observation.

4. Pair-Pair Plots



5. Convergence in delta (the difference between log-likelihood values)

We have used the difference in the log-likelihood function in consecutive iterations for convergence. We stop when the improvement in log-likelihood is less than some threshold value.

The EM for algorithm converges in 25 steps with the stopping condition being $\Delta < 0.1$.




Fig. Convergence of Delta

References

1. https://stephens999.github.io/fiveMinuteStats/intro_to_mixture_models.html
2. https://stephens999.github.io/fiveMinuteStats/intro_to_em.html
3. <http://www.cs.cmu.edu/~guyton/Class/10701-S07/Slides/clustering.pdf>
4. <https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/>

Appendix - Code Snippets

1. Cluster initialization

```
# assigning each data point to a cluster randomly at first
clusters.init <- function(data, k){

  n = nrow(data)
  ## assign each point to random clusters
  clusters = rep(1:k, length.out = n, replace = TRUE)
  data = cbind(data, clusters)

  return(data)
}
```

2. Compute centroid means and variance based on cluster assignment

```
# recalculating the centroid/means and the covariance of all the
clusters, based on the points assigned to them
clusters.parameters <- function(data, k){

  num.features = ncol(data)-1
  #taking mean feature wise
  mean = aggregate(data, by = list(data$clusters), mean, na.rm =
TRUE)
  cov.list <- lapply( sort(unique(data$clusters)), function(x)
cov(data[data$clusters==x,-(num.features+1)],use="na.or.complete"))

  return(list(mean = mean[-1], cov.list = cov.list, cluster_number =
sort(unique(data$clusters)) ))
}
```

3. Compute Euclidean distance and re-assign clusters

```
euc.dist <- function(x1, x2){
  return(sqrt(sum((x1 - x2) ^ 2)))
}
```

```

#re-assign cluster to each data point based on the points euclidean distance from the cluster centroids
assign.cluster.euclidean <- function(data, clusters.list){

  n = nrow(data)
  num.features = ncol(data)-1

  for(i in c(1:n)){
    min_distance = +Inf
    for( j in unique(data$clusters) ){
      mean = clusters.list$mean[j,-(num.features+1)]
      distance = euc.dist(data[i,-ncol(data)], mean )
      if(distance < min_distance){
        min_distance = distance
        #assign the cluster with min distance from the data point
        data[i,]$clusters = j
      }
    }
  }
  return(data)
}

```

4. Main K-means function

```

k.means <- function(data, k){
  data = clusters.init(data, k)
  data.init.clusters = data
  plot.clusters(data)

  ##convergence of dunn index
  prev_dunn_index = 0
  delta = 1

  # stop when dunn index becomes constant
  while(delta!=0){
    print("-----")
    #clusters.list contains the cluster parameters of mean and covariance for each cluster

```

```

clusters.list = clusters.parameters(data, k)
# new column of 'clusters' gets appended
data = assign.cluster.euclidean(data, clusters.list )
plot.clusters(data)

new_dunn_index = dunn_index(data, clusters.list$mean[-6])
print(new_dunn_index)
delta = abs(new_dunn_index - prev_dunn_index)
prev_dunn_index = new_dunn_index
}

return(data)
}

```

5. Dunn Index

A. Intra-cluster Distance for a given cluster

```

intracluster_dist <- function(cluster_mean, cluster_data){
  n = nrow(cluster_data)
  max = 0
  for (i in c(1:n)){
    distance = euc.dist(cluster_mean, cluster_data[i,])
    if(distance>max){
      max = distance
    }
  }
  return(max)
}

```

B. Maximum Intra-cluster Distance for a given cluster

```

get_max_intracluster_distance <- function(data, mean){
  num.features = ncol(data)-1
  for (i in c(1:k)){
    t = lapply(sort(unique(data$clusters)), function(x)
intracluster_dist(mean[x,], data[data$clusters==x,
-(num.features+1)]))
  }
}

```

```
    return(max(unlist(t)))  
  }
```

C. Minimum Intra-cluster Distance between each pair of clusters

```
get_min_intercluster_distance <- function(mean){  
  
  k=3  
  min = +Inf  
  for (i in c(1:(k-1))){  
    for (j in c((i+1):k)){  
      distance = euc.dist(mean[i,], mean[j,])  
      #print(distance)  
      if(distance<min){  
        min = distance  
      }  
    }  
  }  
  return(min)  
}
```

D. Main function for computing Dunn index

```
dunn_index <- function(data, mean){  
  min_intercluster_distance = get_min_intercluster_distance(mean[-6])  
  #print(min_intercluster_distance)  
  max_intracluster_distance = get_max_intracluster_distance(data,  
mean[-6])  
  #print(max_intracluster_distance)  
  return(min_intercluster_distance/ max_intracluster_distance)  
}
```

NOTE - Codes for GMM in the submission zip file.